

Assignment / Explore Query Planning and Indexing

Agnibha Chatterjee

Fall Full 2023

Connecting to database

```
library(RSQLite)
library(DBI)

cwd <- getwd()
dbFile <- "sakila.db"
pathToDb <- paste0(cwd, .Platform$file.sep, dbFile)
conn <- dbConnect(RSQLite::SQLite(), pathToDb)
```

Question 1

Dropping all user-defined indexes

```
sqlStmt <- "DROP INDEX IF EXISTS TitleIndex";
dbSendStatement(conn, sqlStmt)

## <SQLiteResult>
##   SQL  DROP INDEX IF EXISTS TitleIndex
##   ROWS Fetched: 0 [complete]
##       Changed: 0

query <- "SELECT language.name AS language, COUNT(film.film_id) AS no_of_films
        FROM film
        JOIN language ON language.language_id = film.language_id
        GROUP BY language.language_id;
        "
result <- dbGetQuery(conn, query)

## Warning: Closing open result set, pending rows
print(result)

##   language no_of_films
## 1  English          1000
```

Question 2

```
query <- "EXPLAIN QUERY PLAN
        SELECT language.name AS language, COUNT(film.film_id) AS no_of_films
        FROM film
        JOIN language ON language.language_id = film.language_id
        GROUP BY language.name;
```

```

"
result <- dbGetQuery(conn, query)
print(result$detail)

## [1] "SCAN film"
## [2] "SEARCH language USING INTEGER PRIMARY KEY (rowid=?)"
## [3] "USE TEMP B-TREE FOR GROUP BY"

```

Question 3

```

query <- "SELECT film.title as film_name, category.name as category_name, film.length
        FROM film
        JOIN film_category ON film_category.film_id = film.film_id
        AND film.title = 'ZORRO ARK'
        JOIN category ON category.category_id = film_category.category_id;
"

result <- dbGetQuery(conn, query)
print(result)

##   film_name category_name length
## 1 ZORRO ARK      Comedy      50

```

Question 4

```

query <- "EXPLAIN QUERY PLAN
        SELECT film.title as film_name, category.name as category_name, film.length
        FROM film
        JOIN film_category ON film_category.film_id = film.film_id
        AND film.title = 'ZORRO ARK'
        JOIN category ON category.category_id = film_category.category_id;
"

result <- dbGetQuery(conn, query)
print(result$detail)

## [1] "SCAN film_category USING COVERING INDEX sqlite_autoindex_film_category_1"
## [2] "SEARCH category USING INTEGER PRIMARY KEY (rowid=?)"
## [3] "SEARCH film USING INTEGER PRIMARY KEY (rowid=?)"

```

Question 5

```

sqlStmt <- "CREATE INDEX TitleIndex
          ON film (title);
"

result <- dbSendQuery(conn, sqlStmt)
print(result)

## <SQLiteResult>
##   SQL  CREATE INDEX TitleIndex
##         ON film (title);
##
##   ROWS Fetched: 0 [complete]
##         Changed: 0

```

Question 6

```
query <- "EXPLAIN QUERY PLAN
        SELECT film.title as film_name, category.name as category_name, film.length
        FROM film
        JOIN film_category ON film_category.film_id = film.film_id
        AND film.title = 'ZORRO ARK'
        JOIN category ON category.category_id = film_category.category_id;
        "
result <- dbGetQuery(conn, query)

## Warning: Closing open result set, pending rows
print(result$detail)

## [1] "SEARCH film USING INDEX TitleIndex (title=?)"
## [2] "SEARCH film_category USING COVERING INDEX sqlite_autoindex_film_category_1 (film_id=?)"
## [3] "SEARCH category USING INTEGER PRIMARY KEY (rowid=?)"
```

Question 7

The query plan for q6 is different from the one used in q4 because the primary plan displayed above was to use TitleIndex. We can confirm that the index was considered because the output says “SEARCH film USING INDEX TitleIndex”.

Question 8

```
sqlStmt <- "DROP INDEX TitleIndex";
dbSendStatement(conn, sqlStmt)

## <SQLiteResult>
##   SQL  DROP INDEX TitleIndex
##   ROWS Fetched: 0 [complete]
##       Changed: 0

start.time <- Sys.time()
query <- "SELECT film.title as film_name, category.name as category_name, film.length
        FROM film
        JOIN film_category ON film_category.film_id = film.film_id
        AND film.title = 'ZORRO ARK'
        JOIN category ON category.category_id = film_category.category_id;
        "
result <- dbGetQuery(conn, query)

## Warning: Closing open result set, pending rows

end.time <- Sys.time()
time.diff <- end.time - start.time
print(paste("Execution time (without index)", round((time.diff),3), "sec"))

## [1] "Execution time (without index) 0.005 sec"

sqlStmt <- "CREATE INDEX TitleIndex
        ON film (title);
        "
result <- dbSendQuery(conn, sqlStmt)
print(result)
```

```
## <SQLiteResult>
##   SQL  CREATE INDEX TitleIndex
##           ON film (title);
##
##   ROWS Fetched: 0 [complete]
##           Changed: 0

start.time <- Sys.time()
query <- "SELECT f.title as film_name, c.name as category_name, length
        FROM film f
        JOIN film_category fc ON fc.film_id = f.film_id
        AND f.title = 'ZORRO ARK'
        JOIN category c ON c.category_id = fc.category_id;
        "
result <- dbGetQuery(conn, query)

## Warning: Closing open result set, pending rows

end.time <- Sys.time()
time.diff.index <- end.time - start.time
print(paste("Execution time (with index)", round((time.diff),3), "sec"))

## [1] "Execution time (with index) 0.005 sec"

print(paste("Time diff", time.diff - time.diff.index, "sec"))

## [1] "Time diff 0.0029909610748291 sec"
```

We can observe that the time difference between the query without an index and the query with an index is greater than 0, which indicates that fetching the results without an index took longer than fetching them with an index. Although the time difference is relatively small in this case due to the small number of rows in the film table (1,000 rows), the impact of the index would be significantly more noticeable when dealing with a larger table with more number of rows to be scanned.

Question 9

```
query <- "
        SELECT film.title, language.name, film.length
        FROM film
        JOIN language ON language.language_id = film.language_id
        AND LOWER(film.title) LIKE '%gold%';
        "

result <- dbGetQuery(conn, query)
print(result)
```

		title	name	length
## 1	ACE	GOLDFINGER	English	48
## 2	BREAKFAST	GOLDFINGER	English	123
## 3		GOLD RIVER	English	154
## 4	GOLDFINGER	SENSIBILITY	English	93
## 5		GOLDMINE TYCOON	English	153
## 6		OSCAR GOLD	English	115
## 7	SILVERADO	GOLDFINGER	English	74
## 8		SWARM GOLD	English	123

Question 10

```
query <- "  
  EXPLAIN QUERY PLAN  
  SELECT film.title, language.name, film.length  
  FROM film  
  JOIN language ON language.language_id = film.language_id  
  AND LOWER(film.title) LIKE '%gold%';  
  "  
result <- dbGetQuery(conn, query)  
print(result$detail)
```

```
## [1] "SCAN film"  
## [2] "SEARCH language USING INTEGER PRIMARY KEY (rowid=?)"
```

The reason this query doesn't use `TitleIndex` is the use of the `LIKE` operator. The `LIKE` operator's ability to accommodate wildcards introduces complexity for the optimizer in pinpointing the index rows that match the query conditions. To illustrate, consider the query `WHERE title LIKE '%GOLD.'` In this case, an index scan cannot be used because the optimizer cannot determine which index rows commence with the term "GOLD." Similarly, a query like `WHERE title LIKE 'GOLD%'` cannot make use of an index scan because the optimizer cannot identify which index rows conclude with the term "GOLD". A full table scan is performed in this case.

Disconnecting from database

```
dbDisconnect(conn)
```